

NumPy style arrays for C++? [closed]

Asked 8 years, 11 months ago Active 4 months ago Viewed 66k times



91



30



Closed. This question does not meet [Stack Overflow guidelines](#). It is not currently accepting answers.



Want to improve this question? Update the question so it's [on-topic](#) for Stack Overflow.

Closed 6 months ago.

[Improve this question](#)

Are there any C++ (or C) libs that have NumPy-like arrays with support for slicing, vectorized operations, adding and subtracting contents element-by-element, etc.?

c++

arrays

dynamic-arrays

Share Follow

asked Jun 23 '12 at 12:15



[Llamageddon](#)

2,679 ● 2 ● 17 ● 41

2 [Armadillo?](#) – [Oliver Charlesworth](#) Jun 23 '12 at 12:17

1 As far as I know numpy uses [LAPACK](#). While that is written in Fortran, there are c++ bindings available. Never used either of those though. – [Voo](#) Jun 23 '12 at 12:54

Join Stack Overflow to learn, share knowledge, and build your career.

[Sign up with email](#)



[Sign up with Google](#)



[Sign up with GitHub](#)



[Sign up with Facebook](#)



13 Answers

Active

Oldest

Votes



Here are several free software that may suit your needs.

67



1. The [GNU Scientific Library](#) is a GPL software written in C. Thus, it has a C-like allocation and way of programming (pointers, etc.). With the [GSLwrap](#), you can have a C++ way of programming, while still using the GSL. GSL has a [BLAS](#) implementation, but you can use [ATLAS](#) instead of the default CBLAS, if you want even more performances.
2. The [boost/uBLAS](#) library is a BSL library, written in C++ and distributed as a boost package. It is a C++-way of implementing the BLAS standard. uBLAS comes with a few linear algebra functions, and there is an [experimental binding to ATLAS](#).
3. [eigen](#) is a linear algebra library written in C++, distributed under the MPL2 license (starting from version 3.1.1) or LGPL3/GPL2 (older versions). It's a C++ way of programming, but more integrated than the two others (more algorithms and data structures are available). Eigen [claim to be faster](#) than the BLAS implementations above, while not following the de-facto standard BLAS API. Eigen does not seem to put a lot of effort on parallel implementation.
4. [Armadillo](#) is LGPL3 library for C++. It has binding for [LAPACK](#) (the library used by numpy). It uses recursive templates and template meta-programming, which is a good point (I don't know if other libraries are doing it also?).
5. [xtensor](#) is a C++ library that is BSD licensed. It offers A C++ API very similar to that of NumPy. See <https://xtensor.readthedocs.io/en/latest/numpy.html> for a cheat sheet.

These alternatives are really good if you just want to get data structures and basic linear algebra. Depending on your taste about style, license or sysadmin challenges (installing big libraries like LAPACK may be difficult), you may choose the one that best suits your needs.

Share Follow

edited Sep 18 '20 at 15:39



Caio S. Souza

91 ● 1 ● 11

answered Jun 23 '12 at 13:10



nojhan

1,042 ● 9 ● 11

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email



Sign up with Google




Sign up with GitHub



Sign up with Facebook



- 21 Sadly, none of these provide anything as general and convenient as numpy arrays. Numpy arrays are arbitrary-dimensional and support things like `a[:4, ::-1, :, 19] = b[None, -5:, None]` or `a[a>5]=0` and similar, as well as having a huge set of array and index manipulation functions available. I really hope somebody makes something like that for C++ some day. – [amaurea](#) Mar 1 '14 at 11:56
- 4 OpenCV also has a Matrix type that can have arbitrary dimensional size; column/row ranges (`a.colRange(4, 7).rowRange(4, 8)` for `a[4:7, 4, 8]`) and condition mask (`a.setTo(cv::Scalar(0), a>5)` for `a[a>5]=0`) – [xaedes](#) Dec 15 '16 at 16:39
- 3 @amaurea check out the answer on xtensor below, which enables all of the above. – [Quant](#) May 14 '17 at 20:59
- 1 I've had to use Eigen in a recent project and I have to say that while it seems to be efficient, the syntax is absolutely terrible. There is none of that amazing Python slicing syntax available. For example, if you have a 1D vector x and want to manipulate the first n elements, you have to use `x.head(n)`. Don't even ask about manipulating an arbitrary slice of x, you'll need a good old for-loop to do that. This is just one of the many clunky and inconvenient examples I could name. – [Alex](#) May 30 '17 at 23:34 

Try out [xtensor](#). (See the [NumPy to Xtensor Cheat Sheet](#)).

57 xtensor is a C++ library meant for numerical analysis with multi-dimensional array expressions.

xtensor provides



- an extensible expression system enabling numpy-style broadcasting.
- an API following the idioms of the C++ standard library.
- tools to manipulate array expressions and build upon xtensor.

Example

Initialize a 2-D array and compute the sum of one of its rows and a 1-D array.


```
#include <iostream>
#include "xtensor/xarray.hpp"
#include "xtensor/xio.hpp"

xt::xarray<double> arr1
```

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook



```
xt::xarray<double> res = xt::view(arr1, 1) + arr2;

std::cout << res;
```

Outputs

```
{7, 11, 14}
```

Initialize a 1-D array and reshape it inplace.

```
#include <iostream>
#include "xtensor/xarray.hpp"
#include "xtensor/xio.hpp"

xt::xarray<int> arr
    {1, 2, 3, 4, 5, 6, 7, 8, 9};

arr.reshape({3, 3});

std::cout << arr;
```

Outputs

```
{{1, 2, 3},
 {4, 5, 6},
 {7, 8, 9}}
```

Share Follow

edited Dec 3 '17 at 8:39

answered Nov 6 '16 at 11:23





Quant

1,393 ● 13 ● 21

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook



7

all work fine. On the other hand, it is still very experimental and many features haven't been implemented yet.

Here's a simple implementation of the de Casteljau algorithm in C++ using DyND arrays:



```
#include <iostream>
#include <dynd/array.hpp>

using namespace dynd;

nd::array decasteljau(nd::array a, double t){
    size_t e = a.get_dim_size();
    for(size_t i=0; i < e-1; i++){
        a = (1.-t) * a(irange()<(e-i-1)) + t * a(0<irange());
    }
    return a;
}

int main(){
    nd::array a = {1., 2., 2., -1.};
    std::cout << decasteljau(a, .25) << std::endl;
}
```

I wrote a [blog post](#) a little while back with more examples and side-by-side comparisons of the syntax for Fortran 90, DyND in C++, and NumPy in Python.

Disclaimer: I'm one of the current DyND developers.

Share Follow

edited May 19 '16 at 22:14

answered Oct 27 '15 at 17:48



IanH

8,740 ● 1 ● 26 ● 31


3


Eigen is a good linear algebra library.

http://eigen.tuxfamily.org/index.php?title=Main_Page

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook



you need?

Share Follow

answered Jun 23 '12 at 12:20



Frédéric Terrazzoni

1,990 ● 16 ● 25

- 3 The syntax of Eigen is pretty terrible though. There is none of that smooth slicing syntax you find in Numpy. And it's not a general n-dimensional array library, it's more for 1D vectors and 2D matrices only. The fact that they have VectorXd for 1D arrays and MatrixXd for 2D arrays repels me already. – Alex May 30 '17 at 23:50

This is an old question. Still felt like answering. Thought might help many, Especially pydevs coding in C++.

3

If you have already worked with python numpy, then [NumCpp](#) is a great choice. It's minimalistic in syntax and has got similar functions or methods as py numpy.

The [comparison](#) part in the readme doc is also very very cool.



NumCpp

```
nc::NdArray<int> arr = {{4, 2}, {9, 4}, {5, 6}};  
arr.reshape(5, 3);  
arr.astype<double>();
```

Share Follow

answered Oct 12 '20 at 17:22




Sayan Dey


618 ● 1 ● 9


If you want to use multidimensional array(like numpy) for image processing or neural network, you can use `openCV` `cv::Mat` along

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook





vectors with real or complex values, other matrices etc.

A Mat contains the following information: width, height, type, channels, data, flags, datastart, dataend and so on.

It has several methods for matrix manipulation. Bonus you can create then on CUDA cores as well as `cv::cuda::GpuMat`.

Consider I want to create a matrix with 10 rows, 20 columns, type CV_32FC3:

```
int R = 10, C = 20;
Mat m1;
m1.create(R, C, CV_32FC3); //creates empty matrix

Mat m2(cv::Size(R, C), CV_32FC3); // creates a matrix with R rows, C columns
with data type T where R and C are integers,

Mat m3(R, C, CV_32FC3); // same as m2
```

BONUS:

Compile [tiny and compact opencv](#) library for just matrix operations. One of the ways is like as mentioned in this article.


OR


compile opencv source code using following cmake command:

```
$ git clone https://github.com/opencv/opencv.git
$ cd opencv
$ git checkout <version you want to checkout>
$ mkdir build
$ cd build
$ cmake -D WITH_CUDA=OFF -D WITH_MATLAB=OFF -D BUILD_ANDROID_EXAMPLES=OFF -D
BUILD_DOCS=OFF -D BUILD_PERF_TESTS=OFF -D BUILD_TESTS=OFF -
DANDROID_STL=c++_shared -DBUILD_SHARED_LIBS=ON -D BUILD_opencv_objdetect=OFF -D
BUILD_opencv_video=OFF -D BUILD_opencv_videoio=OFF -D
BUILD_opencv_features2d=OFF -D BUILD_opencv_flann=OFF -D
BUILD_opencv_highgui=OFF -D BUILD_opencv_ml=OFF -D BUILD_opencv_photo=OFF -D
```

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook



Try this example:

```
#include "opencv2/core.hpp"
#include<iostream>

int main()
{
    std::cout << "OpenCV Version " << CV_VERSION << std::endl;

    int R = 2, C = 4;
    cv::Mat m1;
    m1.create(R, C, CV_32FC1); //creates empty matrix

    std::cout << "My Mat : \n" << m1 << std::endl;
}
```

Compile the code with following command:

```
$ g++ -std=c++11 opencv_mat.cc -o opencv_mat `pkg-config --libs opencv` `pkg-config --cflags opencv`
```

Run the executable:

```
$ ./opencv_mat

OpenCV Version 3.4.2
My Mat :
[0, 0, 0, 0;
 0, 0, 0, 0]
```

Share Follow

edited Jan 12 at 15:00

answered Jun 18 '20 at 17:11




Milind Deore

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook



2

basic entrywise operations and tensor contractions. Development seems to have slowed down quite some time ago, but perhaps that's just because the library does what it does and not many changes need to be made.



Share Follow

answered May 4 '13 at 13:07



Dan Stahlke

1,349 ● 14 ● 20

xtensor is good, but I ended up writing a mini-library myself as a toy project with c++20, while trying to keep the interface as simple as possible. Here it is: <https://github.com/gbalduzz/NDArray>

2

Example code:



```
using namespace nd;
NDArray<int, 2> m(3, 3); // 3x3 matrix
m = 2; // assign 2 to all
m(-1, all) = 1; // assign 1 to the last row.

auto tile = m(range{1, end}, range{1, end}); // 2x2 tile
std::sort(tile.begin(), tile.end());

std::cout << m; // prints [[2, 2, 2], [2, 1, 1], [1, 2, 2]]
```

~~It does not provide fancy arithmetic operators collapsing multiple operations together, yet,~~ but you can broadcast arbitrary lambdas to a set of tensors with the same shape, or use lazily evaluated arithmetic operators.


Let me know what do you think about the interface and how it compares with the other options, and if this has any hope, what sort of operations you would like to see implemented.

Free license and no dependency!

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email

 Sign up with Google

 Sign up with GitHub

 Sign up with Facebook





Giovanni Balduzzi

21 ● 3



VIGRA contains a good N-dimensional array implementation:

1

<http://ukoethe.github.io/vigra/doc/vigra/Tutorial.html>



I use it extensively, and find it very simple and effective. It's also header only, so very easy to integrate into your development environment. It's the closest thing I've come across to using NumPy in terms of it's API.



The main downside is that it isn't so widely used as the others, so you won't find much help online. That, and it's awkwardly named (try searching for it!)

Share Follow

answered Aug 5 '15 at 14:27



Martin

11 ● 1



Use LibTorch (PyTorch frontend for C++) and be happy.

1

Share Follow



answered Aug 2 '19 at 10:28



Артем Яценко

21 ● 1



[Eigen](#) is a template library for linear algebra (matrices, vectors...). It is header only and free to use (LGPL).

0

Share Follow



answered Jun 23 '12 at 12:21



Claudio

1,570 ● 0 ● 17

Join Stack Overflow to learn, share knowledge, and build your career.

Sign up with email



Sign up with Google



Sign up with GitHub



Sign up with Facebook



0

Share Follow

answered Jun 23 '12 at 12:25



Matt Phillips

8,691 ● 8 ● 41 ● 72



-1



While [GLM](#) is designed to mesh easily with OpenGL and GLSL, it is a fully functional header only math library for C++ with a very intuitive set of interfaces.

It declares vector & matrix types as well as various operations on them.

Multiplying two matrices is as simple as $(M1 * M2)$. Subtracting two vectors $(V1 - V2)$.

Accessing values contained in vectors or matrices is equally simple. After declaring a vec3 vector for example, one can access its first element with `vector.x`. Check it out.

Share Follow

answered Sep 18 '15 at 19:01



user5344731

Join Stack Overflow to learn, share knowledge, and build your career.

[Sign up with email](#)[Sign up with Google](#)[Sign up with GitHub](#)[Sign up with Facebook](#)